

Testing and debugging

Tools for Reproducible Research

Karl Broman

Biostatistics & Medical Informatics, UW–Madison

`kbroman.org`

`github.com/kbroman`

`@kwbroman`

Course web: kbroman.org/Tools4RR

"I tried it, and it worked."

"It's not that we don't test our code, it's that we don't store our tests so they can be re-run automatically."

– Hadley Wickham

R Journal 3(1):5–10, 2011

Types of tests

- ▶ Unit tests
 - For each small function: does it give the right results in specific cases?
- ▶ Integration tests
 - Check that larger multi-function tasks are working.
- ▶ Regression tests
 - Compare output to saved results, to check that things that worked continue working.

Types of tests

- ▶ Check inputs
 - Stop if the inputs aren't as expected.
- ▶ Unit tests
 - For each small function: does it give the right results in specific cases?
- ▶ Integration tests
 - Check that larger multi-function tasks are working.
- ▶ Regression tests
 - Compare output to saved results, to check that things that worked continue working.

Check inputs

```
winsorize <-  
function(x, q=0.006)  
{  
  if(!is.numeric(x)) stop("x should be numeric")  
  
  if(!is.numeric(q)) stop("q should be numeric")  
  if(length(q) > 1) {  
    q <- q[1]  
    warning("length(q) > 1; using q[1]")  
  }  
  if(q < 0 || q > 1) stop("q should be in [0,1]")  
  
  lohi <- quantile(x, c(q, 1-q), na.rm=TRUE)  
  if(diff(lohi) < 0) lohi <- rev(lohi)  
  
  x[!is.na(x) & x < lohi[1]] <- lohi[1]  
  x[!is.na(x) & x > lohi[2]] <- lohi[2]  
  x  
}
```

Check inputs

```
winsorize <-  
function(x, q=0.006)  
{  
  stopifnot(is.numeric(x))  
  stopifnot(is.numeric(q), length(q)==1, q>=0, q<=1)  
  
  lohi <- quantile(x, c(q, 1-q), na.rm=TRUE)  
  if(diff(lohi) < 0) lohi <- rev(lohi)  
  
  x[!is.na(x) & x < lohi[1]] <- lohi[1]  
  x[!is.na(x) & x > lohi[2]] <- lohi[2]  
  x  
}
```

assertthat package

```
#' import assertthat
winsorize <-
function(x, q=0.006)
{
  if(all(is.na(x)) || is.null(x)) return(x)

  assert_that(is.numeric(x))
  assert_that(is.number(q), q>=0, q<=1)

  lohi <- quantile(x, c(q, 1-q), na.rm=TRUE)
  if(diff(lohi) < 0) lohi <- rev(lohi)

  x[!is.na(x) & x < lohi[1]] <- lohi[1]
  x[!is.na(x) & x > lohi[2]] <- lohi[2]
  x
}
```


Tests in R packages

- ▶ Examples in .Rd files
- ▶ Vignettes
- ▶ tests/ directory
 - `some_test.R` and `some_test.Rout.save`

R CMD check is your friend.

An example example

```
#' @examples
#' x <- sample(c(1:10, rep(NA, 10), 21:30))
#' winsorize(x, 0.2)
```

A tests/ example

```
library(qtl)

# read data
csv <- read.cross("csv", "", "listeria.csv")

# write
write.cross(csv, "csv", filestem="junk")

# read back in
csv2 <- read.cross("csv", "", "junk.csv",
                  genotypes=c("AA", "AB", "BB",
                              "not BB", "not AA"))

# check for a change
comparecrosses(csv, csv2)

unlink("junk.csv")
```

testthat package

- ▶ Expectations

```
expect_equal(10, 10 + 1e-7)
expect_identical(10, 10)
expect_equivalent(c("one"=1), 1)
expect_warning(log(-1))
expect_error(1 + "a")
```

- ▶ Tests

```
test_that("winsorize small vectors", { ... })
```

- ▶ Contexts

```
context("Group of related tests")
```

- ▶ Store tests in inst/tests/

- ▶ tests/run-all.R file containing

```
library(testthat)
library(mypkg)
test_package("mypkg")
```

Example testthat test

```
context("winsorize")

test_that("winsorize works for small vectors", {

  x <-      c(2, 3, 7, 9, 6, NA, 5, 8, NA, 0, 4, 1, 10)
  result1 <- c(2, 3, 7, 9, 6, NA, 5, 8, NA, 1, 4, 1, 9)
  result2 <- c(2, 3, 7, 8, 6, NA, 5, 8, NA, 2, 4, 2, 8)

  expect_identical(winsorize(x, 0.1), result1)
  expect_identical(winsorize(x, 0.2), result2)

})
```

Example testthat test

```
test_that("winsorize works for a long vector", {  
  
  set.seed(94745689)  
  n <- 1000  
  nmis <- 10  
  p <- 0.05  
  input <- rnorm(n)  
  input[sample(1:n, nmis)] <- NA  
  quL <- quantile(input, p, na.rm=TRUE)  
  quH <- quantile(input, 1-p, na.rm=TRUE)  
  
  result <- winsorize(input, p)  
  middle <- !is.na(input) & input >= quL & input <= quH  
  low <- !is.na(input) & input <= quL  
  high <- !is.na(input) & input >= quH  
  
  expect_identical(is.na(input), is.na(result))  
  expect_identical(input[middle], result[middle])  
  expect_true( all(result[low] == quL) )  
  expect_true( all(result[high] == quH) )  
  
})
```

Workflow

- ▶ Write tests as you're coding.
- ▶ Run `test()`
with `devtools`, and working in your package directory
- ▶ Consider `auto_test("R", "tests")`
automatically runs tests when any file changes
- ▶ Periodically run `R CMD check`
also `R CMD check --as-cran`

What to test?

- ▶ You can't test **everything**.
- ▶ Focus on the **boundaries**
 - (Depends on the nature of the problem)
 - Vectors of length 0 or 1
 - Things exactly matching
 - Things with no matches
- ▶ Test handling of missing data.
NA, Inf, -Inf
- ▶ Automate the construction of test cases
 - Create a table of inputs and expected outputs
 - Run through the values in the table

Another example

```
test_that("running mean with constant x or position", {  
  
  n <- 100  
  x <- rnorm(n)  
  pos <- rep(0, n)  
  
  expect_equal( runningmean(pos, x, window=1), rep(mean(x), n) )  
  expect_equal( runningmean(pos, x, window=1, what="median"),  
                rep(median(x), n) )  
  expect_equal( runningmean(pos, x, window=1, what="sd"),  
                rep(sd(x), n) )  
  
  x <- rep(0, n)  
  pos <- runif(n, 0, 5)  
  
  expect_equal( runningmean(pos, x, window=1), x)  
  expect_equal( runningmean(pos, x, window=1, what="median"), x)  
  expect_equal( runningmean(pos, x, window=5, what="sd"),  
                rep(0, n))  
})
```

Debugging tools

- ▶ `cat`, `print`
- ▶ `traceback`, `browser`, `debug`
- ▶ RStudio breakpoints
- ▶ Eclipse/StatET
- ▶ `gdb`

Debugging

Step 1: Reproduce the problem

Debugging

Step 1: Reproduce the problem

Step 2: Turn it into a test

Debugging

Isolate the problem: where do things go bad?

Debugging

Don't make the same mistake twice.

The most pernicious bugs

The code is right, but your thinking is wrong.

The most pernicious bugs

The code is right, but your thinking is wrong.

You were mistaken about what the code would do.

The most pernicious bugs

The code is right, but your thinking is wrong.

You were mistaken about what the code would do.

→ Write trivial programs to test your understanding.

Summary

- ▶ If you don't test your code, how do you know it works?
- ▶ If you test your code, save and automate those tests.
- ▶ Check the input to each function.
- ▶ Write unit tests for each function.
- ▶ Write some larger regression tests.
- ▶ Turn bugs into tests.